

IBM Data Engineering Capstone Project

Project Overview

As part of the Capstone project, you will assume the role of an Associate Data Engineer who has recently joined an e-commerce organization. You will be presented with a business challenge that requires building a data platform for retailer data analytics.

In this Capstone project, you will:

Design a data platform that uses MySQL as an OLTP database and MongoDB as a NoSQL database.

Design and implement a data warehouse and generate reports from the data.

Design a reporting dashboard that reflects the key metrics of the business.

Extract data from OLTP and NoSQL databases, transform it and load it into the data warehouse, and then create an ETL pipeline.

And finally, create a Spark connection to the data warehouse, and then deploy a machine learning model.

In Module 1, you will design the OLTP database for an E-Commerce website, populate the OLTP Database with the data provided and automate the export of the daily incremental data into the data warehouse.

In Module 2, you will set up a NoSQL database to store the catalogue data for an E-Commerce website, load the E-Commerce catalogue data into the NoSQL database, and query the E-Commerce catalogue data in the NoSQL database.

In Module 3, you will design the schema for a data warehouse based on the schema of the OLTP and NoSQL databases. You'll then create the schema and load the data into fact and dimension tables, automate the daily incremental data insertion into the data warehouse, and create Cubes and Rollups to make the reporting easier.

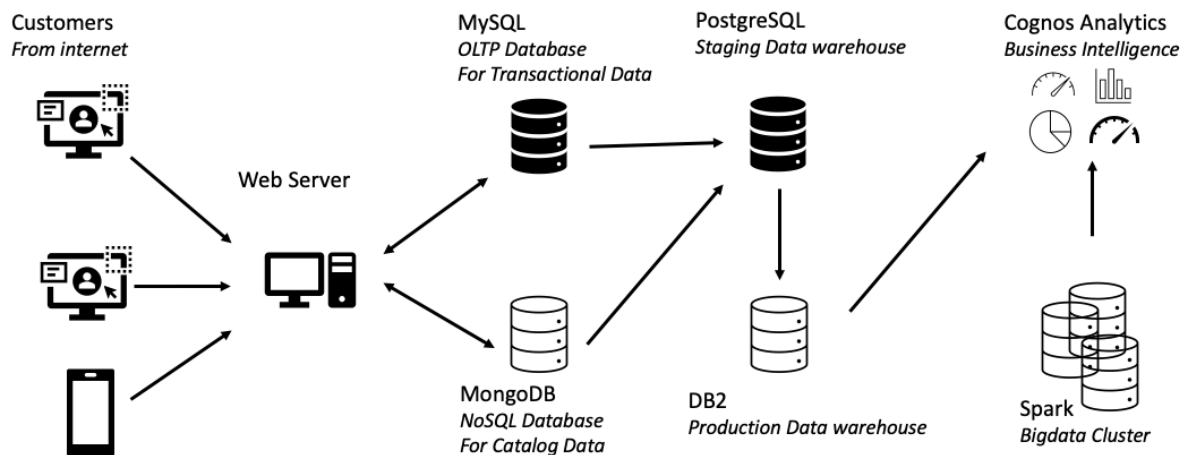
In Module 4, you will create a Cognos data source that points to a data warehouse table, create a bar chart of Quarterly sales of cell phones, create a pie chart of sales of electronic goods by category, and create a line chart of total sales per month for the year 2020.

In Module 5, you will extract data from OLTP, NoSQL, and MongoDB databases into CSV format. You will then transform the OLTP data to suit the data warehouse schema and then load the transformed data into the data warehouse. Finally, you will verify that the data is loaded properly.

In the sixth and final module, you will use your skills in Big Data Analytics to create a Spark connection to the data warehouse, and then deploy a machine learning model on SparkML for making sales projections.

Data Platform Architecture

Data Platform Architecture



This is the data platform architecture of an ecommerce company named SoftCart. SoftCart uses a hybrid architecture, with some of its databases on premises and some on cloud.

Tools and Technologies:

- OLTP database - MySQL
- NoSQL database - MongoDB
- Production Data warehouse – DB2 on Cloud
- Staging - Data warehouse – PostgreSQL
- Big data platform - Hadoop
- Big data analytics platform – Spark
- Business Intelligence Dashboard - IBM Cognos Analytics
- Data Pipelines - Apache Airflow

Process:

SoftCart's online presence is primarily through its website, which customers access using a variety of devices like laptops, mobiles and tablets. All the catalogue data of the products is stored in the MongoDB NoSQL server. All the transactional data like inventory and sales are stored in the MySQL database server. SoftCart's webserver is driven entirely by these two databases. Data is periodically extracted from these two databases and put into the staging data warehouse running on PostgreSQL. Production data warehouse is on the cloud instance of IBM DB2 server. BI teams connect to the IBM DB2 for operational dashboard creation. IBM Cognos Analytics is used to create dashboards. SoftCart uses Hadoop cluster as its big data platform where all the data collected for analytics purposes. Spark is used to analyse the data on the Hadoop cluster. To move data between OLTP, NoSQL and the data warehouse, ETL pipelines are used and these run on Apache Airflow.

Module 1: with link to OLTP database exercise

OLTP database requirements and design

OLTP database:

OLTP database is generally used to handle everyday business transactions of an organization like a bank or a super market chain. OLTP databases can be write heavy or may have a balanced read/write load.

OLTP database requirements:

An OLTP database is expected to handle a huge number of transactions per second. Each transaction usually involves accessing (read/write) a small portion of the database, in other words the payload per transaction is small.

The time taken to execute a transaction usually called latency needs to be very less.

OLTP database design:

The schema of an OLTP database is highly normalized so as to achieve a very low latency. To further improve the latency an OLTP database stores only the recent data like the last few week's data. They are usually run on storage that is very fast like SSD.

Scenario

You are a data engineer at an e-commerce company. Your company needs you to design a data platform that uses MySQL as an OLTP database. You will be using MySQL to store the OLTP data.

Design the OLTP Database

Creating sales database and sales_data table in sales database.

```
Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

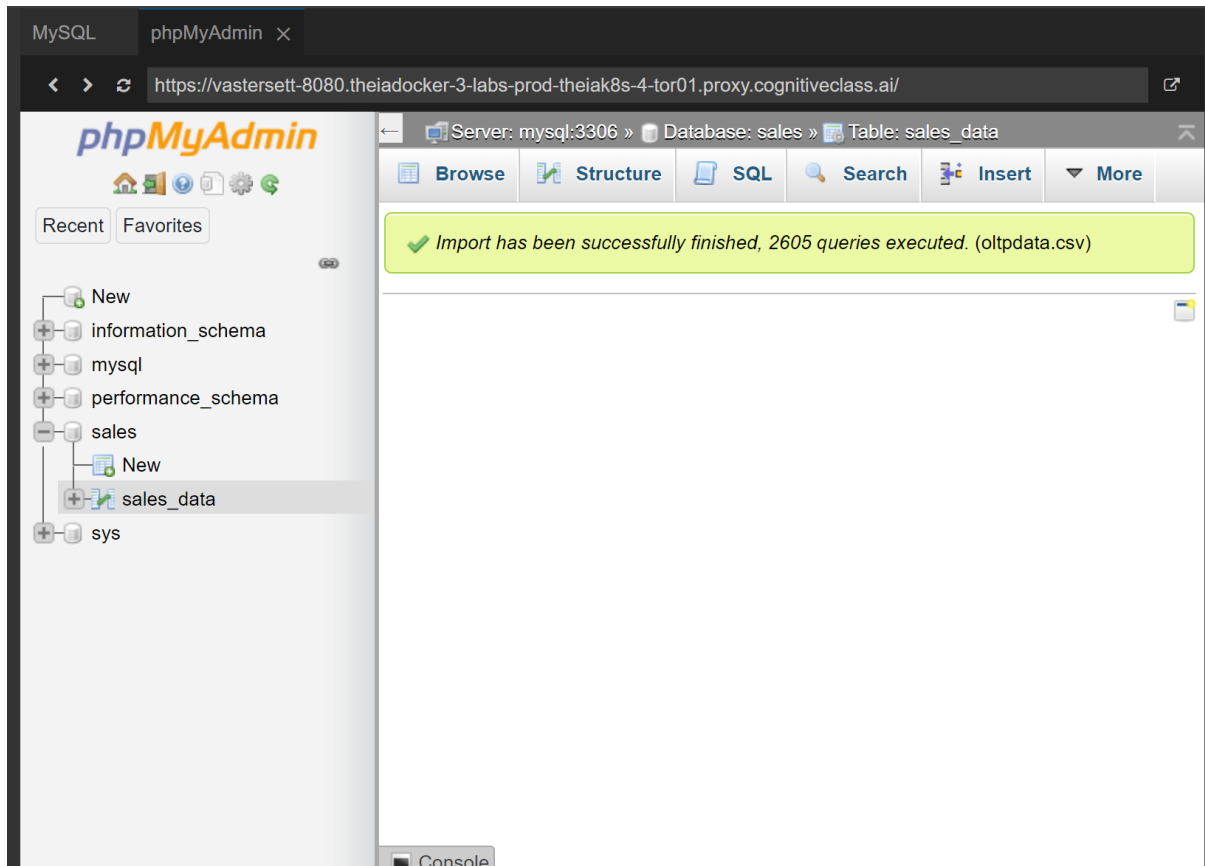
mysql> CREATE DATABASE sales;
Query OK, 1 row affected (0.01 sec)

mysql> USE sales;
Database changed
mysql> CREATE TABLE sales_data (
    -> product_id INT,
    -> customer_id INT,
    -> price INT,
    -> quantity INT,
    -> timestamp DATETIME);
Query OK, 0 rows affected (0.04 sec)

mysql> 
```

Load the given data(oltpdata.csv) into sales_data table

Importing data from oltpdata.csv file into sales_data table using phpMyAdmin GUI.



List the tables in the sales database and query the count of records in the sales_data table to make sure the data are imported correctly and successfully.

```
mysql> SHOW TABLES;
+-----+
| Tables_in_sales |
+-----+
| sales_data      |
+-----+
1 row in set (0.00 sec)

mysql> 
```

```
mysql> SELECT COUNT(*) AS Records_in_sales_data FROM sales_data;
+-----+
| Records_in_sales_data |
+-----+
|                2605 |
+-----+
1 row in set (0.01 sec)

mysql> 
```

Set up Admin tasks

Creating index on the timestamp field to speed up queries.

```
mysql> CREATE INDEX ts ON sales_data(timestamp);
Query OK, 0 rows affected (0.09 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> SHOW INDEX FROM sales_data;
+-----+-----+-----+-----+-----+-----+-----+-----+
| Table | Non_unique | Key_name | Seq_in_index | Column_name | Collation | Cardinality | Sub_part |
| Packed | Null | Index_type | Comment | Index_comment | Visible | Expression |
+-----+-----+-----+-----+-----+-----+-----+
| sales_data | 1 | ts | 1 | timestamp | A | 2605 | NULL |
| NULL | YES | BTREE | | | YES | NULL |
+-----+-----+-----+-----+-----+-----+
1 row in set (0.04 sec)

mysql>
```

Writing a bash script(datadump.sh) that exports all the rows in the sales_data table to a file named sales_data.sql to automate the export of the daily incremental data into the data warehouse.

```
MySQL      phpMyAdmin      datadump.sh X
datadump.sh
1  #!/bin/sh
2  # The above line tells the interpreter this code needs to be run as a shell script.
3
4  # This will be printed on to the screen. In the case of cron job, it will be printed to the logs.
5  echo "Pulling Database: This may take a few minutes"
6  # Create a backup
7  if mysqldump --user=root --password='NzQzNS12YXN0ZXJz' sales sales_data > sales_data.sql ; then
8  | echo 'sales_data.sql created'
9  | else
10 | echo 'Error sales_data.sql was not created!'
11 | exit
12 | fi
```

```
Problems      theia@theiadocker-vastersett: /home/project      theia@theiadocker-vastersett: /home/project X

theia@theiadocker-vastersett:/home/project$ ls -l datadump.sh
-rw-r--r-- 1 theia users 450 Jul  8 05:29 datadump.sh
theia@theiadocker-vastersett:/home/project$ chmod +x datadump.sh
theia@theiadocker-vastersett:/home/project$ ls -l datadump.sh
-rwxr-xr-x 1 theia users 450 Jul  8 05:29 datadump.sh
theia@theiadocker-vastersett:/home/project$ ./datadump.sh
Pulling Database: This may take a few minutes
mysqldump: [Warning] Using a password on the command line interface can be insecure.
sales_data.sql created
theia@theiadocker-vastersett:/home/project$ ls -l
total 220
-rwxr-xr-x 1 theia users 450 Jul  8 05:29 datadump.sh
-rw-r--r-- 1 theia users 218964 Jul  8 05:34 sales_data.sql
theia@theiadocker-vastersett:/home/project$
```

Module 2:

Scenario

You are a data engineer at an e-commerce company. Your company needs you to design a data platform that uses MongoDB as a NoSQL database. You will be using MongoDB to store the e-commerce catalogue data.

Importing data into MongoDB database

Installing mongoimport and mongoexport. Importing given data(catalog.json) into a database named 'catalog' and a collection named 'electronics' on mongodb server.

```
Problems theia@theiadocker-vastersett: /home/project X
theia@theiadocker-vastersett:/home/project$ mongoimport --version
bash: mongoimport: command not found
theia@theiadocker-vastersett:/home/project$ wget https://fastdl.mongodb.org/tools/db/mongodb-database-tools-ubuntu1804-x86_64-100.3.1.tgz
--2022-07-11 12:09:08-- https://fastdl.mongodb.org/tools/db/mongodb-database-tools-ubuntu1804-x86_64-100.3.1.tgz
Resolving fastdl.mongodb.org (fastdl.mongodb.org)... 54.192.51.20, 54.192.51.124, 54.192.51.66, ...
Connecting to fastdl.mongodb.org (fastdl.mongodb.org)|54.192.51.20|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 64021650 (61M) [binary/octet-stream]
Saving to: 'mongodb-database-tools-ubuntu1804-x86_64-100.3.1.tgz'

mongodb-database-tools-u 100%[=====>] 61.06M 33.3MB/s in 1.8s

2022-07-11 12:09:11 (33.3 MB/s) - 'mongodb-database-tools-ubuntu1804-x86_64-100.3.1.tgz' saved [64021650/64021650]

theia@theiadocker-vastersett:/home/project$ tar -xf mongodb-database-tools-ubuntu1804-x86_64-100.3.1.tgz
theia@theiadocker-vastersett:/home/project$ export PATH=$PATH:/home/project/mongodb-database-tools-ubuntu1804-x86_64-100.3.1/bin
theia@theiadocker-vastersett:/home/project$ echo "done"
done
theia@theiadocker-vastersett:/home/project$ mongoimport --version
mongoimport version: 100.3.1
git version: 32632b931f9c41d8314b75ecc88e551b012b1e30
Go version: go1.15.8
os: linux
arch: amd64
compiler: gc
theia@theiadocker-vastersett:/home/project$ mongoimport -u root -p MTMwMzUtdmFzdGVy --authenticationDatabase admin --db catalog --collection electronics --file catalog.json
2022-07-11T12:10:02.320+0000 connected to: mongodb://localhost/
2022-07-11T12:10:02.357+0000 438 document(s) imported successfully. 0 document(s) failed to import.
theia@theiadocker-vastersett:/home/project$
```

List out all databases and collections on mongodb server to check if catalog database and electronics collection are created successfully or not.

```
theia@theiadocker-vastersett:/home/project$ mongo -u root -p MTMwMzUtdmFzdGVy --authenticationDatabase admin local
MongoDB shell version v3.6.3
connecting to: mongodb://127.0.0.1:27017/local
MongoDB server version: 3.6.3
Server has startup warnings:
2022-07-11T12:01:11.984+0000 I STORAGE [initandlisten]
2022-07-11T12:01:11.984+0000 I STORAGE [initandlisten] ** WARNING: Using the XFS filesystem is strongly recommended with the WiredTiger storage engine
2022-07-11T12:01:11.984+0000 I STORAGE [initandlisten] ** See http://dochub.mongodb.org/core/prodnotes-filesystem
> show dbs
admin      0.000GB
catalog    0.000GB
config     0.000GB
local      0.000GB
> █
```

```
> use catalog
switched to db catalog
> show collections
electronics
> █
```

Creating index on the field 'type' for faster queries.

```
> db.electronics.createIndex({"type":1})
{
  "createdCollectionAutomatically" : false,
  "numIndexesBefore" : 1,
  "numIndexesAfter" : 2,
  "ok" : 1
}
```

Trying out a few queries to check the imported data:

Finding the count of laptops.

```
> db.electronics.count({"type":"laptop"})
389
> █
```

Finding the number of smart phones with screen size of 6 inches.

```
> db.electronics.count({"type":"smart phone","screen size":6})
8
> █
```

Finding the average screen size of smart phones.

```
> db.electronics.aggregate([{"$group":{"_id":"$type","average":{"$avg":"$screen size"}}}]
{ "_id" : "television", "average" : 39.8 }
{ "_id" : "laptop", "average" : 14.568123393316196 }
{ "_id" : "smart phone", "average" : 6 }
> db.electronics.aggregate([{"$match":{"type":"smart phone"}},{"$group":{"_id":"$type","average":{"$avg":"$screen size"}}}]
{ "_id" : "smart phone", "average" : 6 }
> █
```

Export data from MongoDB database as csv file

Exporting the fields `_id`, “type”, “model”, from the ‘electronics’ collection into a file named `electronics.csv`

```
> exit
bye
theia@theiadocker-vastersett:/home/project$ mongoexport -u root -p MTMwMzUtdmFzdGVy --authenticationDatabase admin --db catalog --collection electronics --out electronics.csv --type=csv --fields _id,type,model
2022-07-11T12:49:29.772+0000    connected to: mongodb://localhost/
2022-07-11T12:49:29.786+0000    exported 438 records
theia@theiadocker-vastersett:/home/project$
```

Module 3.1:

Scenario

You are a data engineer hired by an ecommerce company named SoftCart.com. The company retails download only items like E-Books, Movies, Songs etc. The company has international presence and customers from all over the world. The company would like to create a data warehouse so that it can create reports like:

- total sales per year per country
- total sales per month per category
- total sales per quarter per country
- total sales per category per country

You will use your data warehousing skills to design and implement a data warehouse for the company.

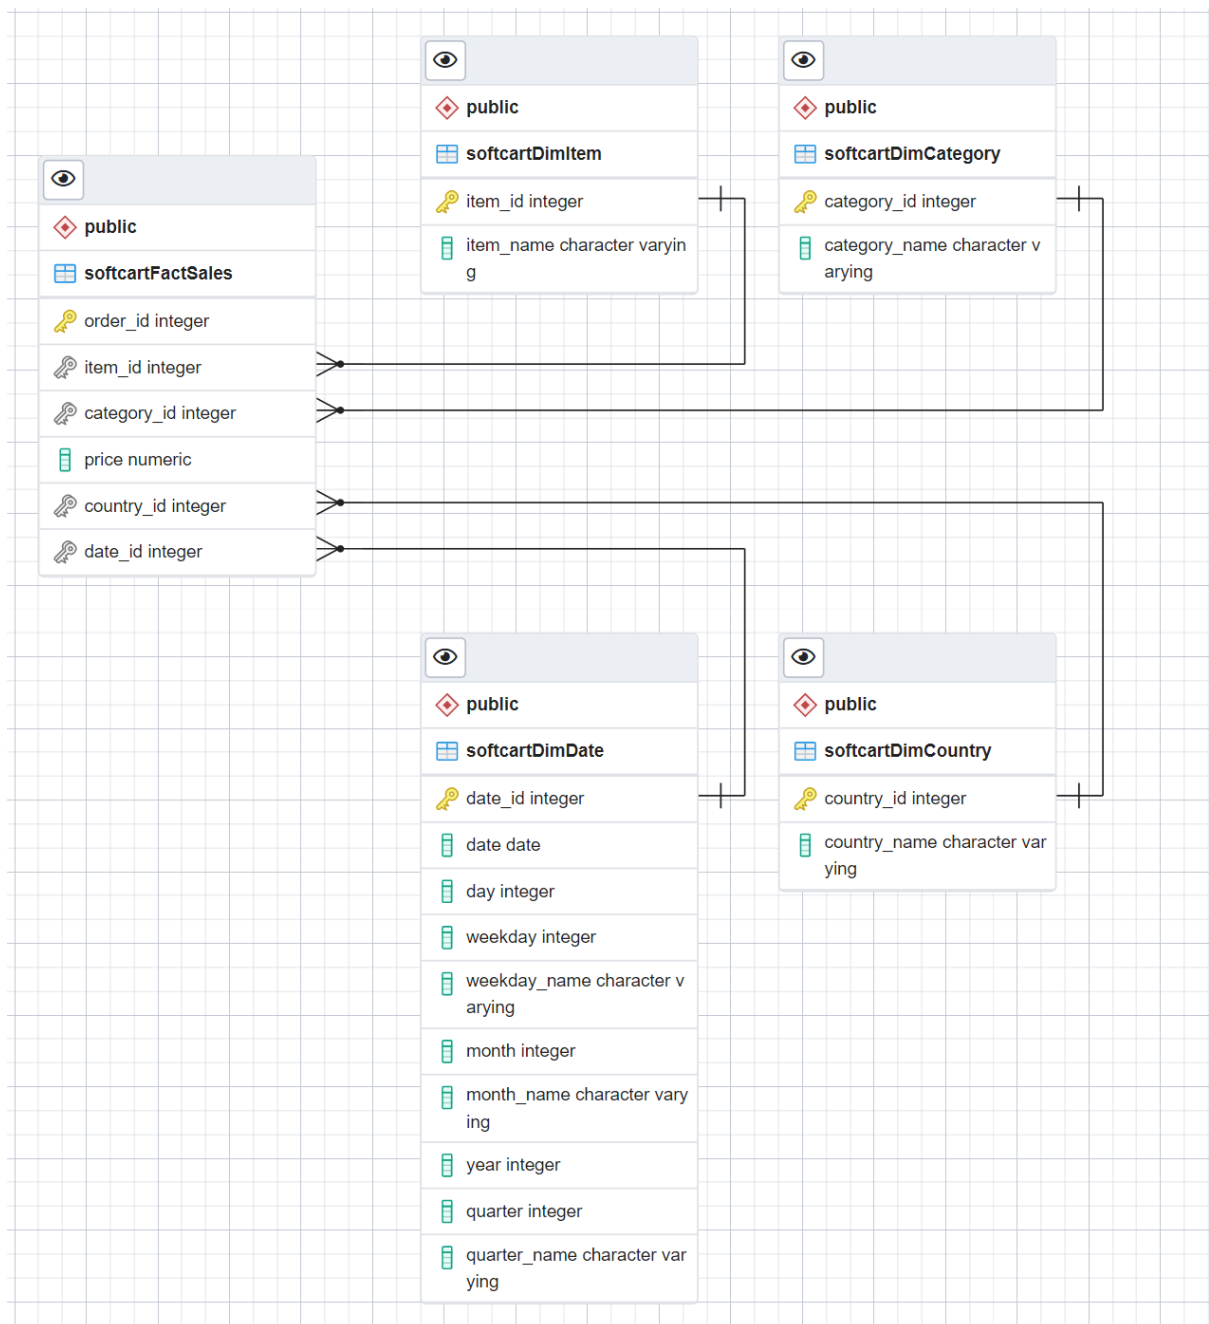
Designing a Data Warehouse

The ecommerce company has provided you the sample data.

OrderID	Item	Category	Price	Country	Date
2123	The Matrix	Movie	9.99	USA	20-Feb-21
3254	The Alchemist	Ebook	5.99	Canada	20-Feb-21
4901	Baby Shark	Song	2.49	Japan	20-Feb-21
5679	The Lord of the Rings	Ebook	6.99	Cyprus	20-Feb-21

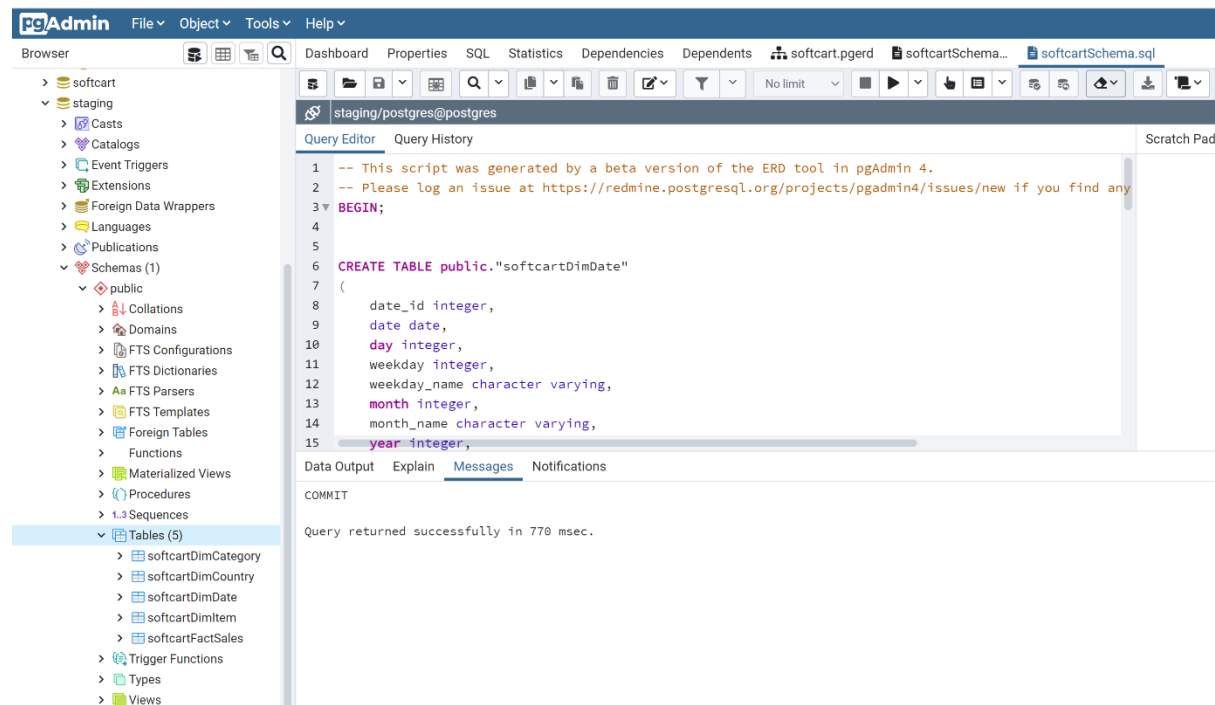
You will start your project by designing a Star Schema for the warehouse by identifying the columns for the various dimension and fact tables in the schema. Name your database as `softcart`.

Creating tables and relationships amongst created tables, using the ERD design Tool of pgAdmin.



Create the schema

After creating tables and relationships using ERD design tool, generate the schema sql using the tool. Then use the created schema sql to create the schema in a database named staging.



Module 3.2:

Scenario

You are a data engineer hired by an ecommerce company named SoftCart.com. The company retails download only items like E-Books, Movies, Songs etc. The company has international presence and customers from all over the world. You have designed the schema for the data warehouse in the previous assignment. Data engineering is a team game. Your senior data engineer reviewed your design. Your schema design was improvised to suit the production needs of the company. In this assignment you will generate reports out of the data in the data warehouse.

Load data into the Data Warehouse (IBM DB2)

Loading the data provided by the company in csv format into the tables (DimDate, DimCategory, DimCountry, FactSales).

IBM Db2 on Cloud

Load Data Load History Tables Views Indexes Aliases MQTs Sequences Application objects

HMM70408.DIMDATE

Back

Export to CSV

	DATEID SMALLINT	DATE DATE	YEAR SMALLINT	QUARTER SMALLINT	QUARTERNAME VARCHAR(2)	MONTH SMALLINT	MONTHNAME VARCHAR(9)	DAY SMALLINT	WEEKDAY SMALLINT	WEEKDAYN VARCHAR(9)
1	1	2019-01-01	2019	1	Q1	1	January	1	3	Tuesday
2	2	2019-01-02	2019	1	Q1	1	January	2	4	Wednesday
3	3	2019-01-03	2019	1	Q1	1	January	3	5	Thursday
4	4	2019-01-04	2019	1	Q1	1	January	4	6	Friday
5	5	2019-01-05	2019	1	Q1	1	January	5	7	Saturday
6	6	2019-01-06	2019	1	Q1	1	January	6	1	Sunday
7	7	2019-01-07	2019	1	Q1	1	January	7	2	Monday
8	8	2019-01-08	2019	1	Q1	1	January	8	3	Tuesday
9	9	2019-01-09	2019	1	Q1	1	January	9	4	Wednesday
10	10	2019-01-10	2019	1	Q1	1	January	10	5	Thursday
11	11	2019-01-11	2019	1	Q1	1	January	11	6	Friday
12	12	2019-01-12	2019	1	Q1	1	January	12	7	Saturday
13	13	2019-01-13	2019	1	Q1	1	January	13	1	Sunday
14	14	2019-01-14	2019	1	Q1	1	January	14	2	Monday
15	15	2019-01-15	2019	1	Q1	1	January	15	3	Tuesday
16	16	2019-01-16	2019	1	Q1	1	January	16	4	Wednesday

IBM Db2 on Cloud

Load Data Load History Tables Views Indexes Aliases MQTs Sequences Application objects

HMM70408.DIMCATEGORY

Back

Export to CSV

	CATEGORYID SMALLINT	CATEGORY VARCHAR(11)
1	1	Electronics
2	2	Books
3	3	Toys
4	4	Sports
5	5	Software

IBM Db2 on Cloud

Load DataLoad HistoryTablesViewsIndexesAliasesMQTSSequencesApplication objects

HHM70408.DIMCOUNTRY

Back

Export to CSV

	COUNTRYID SMALLINT	COUNTRY VARCHAR(20)
1	1	Argentina
2	2	Australia
3	3	Austria
4	4	Azerbaijan
5	5	Belgium
6	6	Brazil
7	7	Bulgaria
8	8	Canada
9	9	Cyprus
10	10	Czech Republic
11	11	Denmark
12	12	Egypt
13	13	Estonia
14	14	Finland
15	15	France
16	16	Germany

IBM Db2 on Cloud

Load Data Load History Tables Views Indexes Aliases MQTs Sequences Application objects

HMM70408.FACTSALES Back

SQL

Export to CSV

	ORDERID INTEGER	DATEID SMALLINT	COUNTRYID SMALLINT	CATEGORYID SMALLINT	AMOUNT SMALLINT
1	1	1	6	4	5190
2	2	1	25	2	1205
3	3	1	22	5	3155
4	4	1	40	2	268
5	5	1	28	3	3199
6	6	1	30	4	6643
7	7	1	11	2	3995
8	8	1	14	4	2318
9	9	1	18	2	3956
10	10	1	37	1	3681
11	11	1	10	3	1309
12	12	1	1	1	5937
13	13	1	9	2	6487
14	14	1	42	1	7745
15	15	1	6	1	3597
16	16	1	11	5	399

Queries for data analytics

Creating and running queries to check if the created data warehouse can create reports.

Grouping sets query using the columns country, category, totalsales. (Total sales per country and Total sales per category)

The screenshot shows the IBM Db2 on Cloud interface with a SQL query editor and a results pane. The query is a GROUP BY GROUPING SETS query that calculates total sales by country and category.

```
1 SELECT country, category, SUM(amount) AS "Total_Sales"
2 FROM factsales
3 LEFT JOIN dimcountry
4 ON factsales.countryid = dimcountry.countryid
5 LEFT JOIN dimcategory
6 ON factsales.categoryid = dimcategory.categoryid
7 GROUP BY GROUPING SETS(country, category)
8 ORDER BY country, category
```

The results pane displays the following data:

COUNTRY	CATEGORY	Total_Sales
Argentina		21755581
Australia		21522004
Austria		21365726
Azerbaijan		21325766
Belgium		21498249
	Books	239357597
	Electronics	239912568
	Software	240289802

Rollup query using the columns year, country, and totalsales. (Total sales per year per country and Total sales per year)

The screenshot shows the IBM Db2 on Cloud interface with a SQL query editor and a results pane. The query is a ROLLUP query that calculates total sales by year and country.

```
1 SELECT year, country, SUM(amount) AS "Total_Sales"
2 FROM factsales
3 LEFT JOIN dimdate
4 ON factsales.dateid = dimdate.dateid
5 LEFT JOIN dimcountry
6 ON factsales.countryid = dimcountry.countryid
7 GROUP BY ROLLUP(year, country)
8 ORDER BY year, country
```

The results pane displays the following data:

YEAR	COUNTRY	Total_Sales
2019	Vietnam	7064760
2019		399729036
2020	Argentina	7327655
2020	Australia	6964260

Cube query using the columns year, country, and average sales. (Average sales per year per country, Average sales per year, Average sales per country)

The screenshot shows the IBM Db2 on Cloud Run SQL interface. The SQL editor contains a query using the CUBE operator to calculate average sales per year and country. The query is as follows:

```

1 SELECT year, country, AVG(amount) AS "Average_Sales"
2 FROM factsales
3 LEFT JOIN dimdate
4 ON factsales.dateid = dimdate.dateid
5 LEFT JOIN dimcountry
6 ON factsales.countryid = dimcountry.countryid
7 GROUP BY CUBE(year, country)
8 ORDER BY year, country

```

The results are displayed in two windows. The main window shows the first 5 rows of the result set:

YEAR	COUNTRY	Average_Sales
2019	Argentina	4017
2019	Australia	4066
2019	Austria	4078
2019	Azerbaijan	3967
2019	Belgium	3978

A smaller window also displays the first 5 rows of the result set:

YEAR	COUNTRY	Average_Sales
2019	Vietnam	3946
2019		3999
2020	Argentina	4114
2020	Australia	3899

Creating a materialized query table (MQT) named total_sales_per_country that has the columns country and total_sales to improve the performance of complex queries that operate on very large amounts of data.

Db2 uses a materialized query table to precompute the results of data that is derived from one or more tables. When you submit a query, Db2 can use the results that are stored in a materialized query table rather than compute the results from the underlying source tables on which the materialized query table is defined.

The screenshot shows the IBM Db2 on Cloud Run SQL interface. The SQL editor contains a query to create a materialized query table (MQT) named total_sales_per_country. The query is as follows:

```

1 CREATE TABLE Total_Sales_per_Country (country, Total_Sales) AS
2 (SELECT country, SUM(amount) AS "Total_Sales"
3 FROM factsales
4 LEFT JOIN dimcountry
5 ON factsales.countryid = dimcountry.countryid
6 GROUP BY country)
7 DATA INITIALLY DEFERRED
8 REFRESH DEFERRED
9 MAINTAINED BY SYSTEM;
10 REFRESH TABLE total_sales_per_country;
11 SELECT * FROM total_sales_per_country;

```

The results are displayed in two windows. The main window shows the results of the CREATE TABLE statement, which includes a warning message:

Warning message
The materialized query table "HHM70408.TOTAL_SALES_PER_COUNTRY" may not be used to optimize the processing of queries... SQLCODE=20059, SQLSTATE=01633, DRIVER=4.31.10

The results of the REFRESH TABLE statement are also shown:

COUNTRY	TOTAL_SALES
Czech Republic	21334142
Denmark	21331097
Egypt	21379967
Estonia	21493054

A smaller window also displays the results of the SELECT statement:

COUNTRY	TOTAL_SALES
Argentina	21755581
Australia	21522004
Austria	21365726
Azerbaijan	21325766
Belgium	21498249

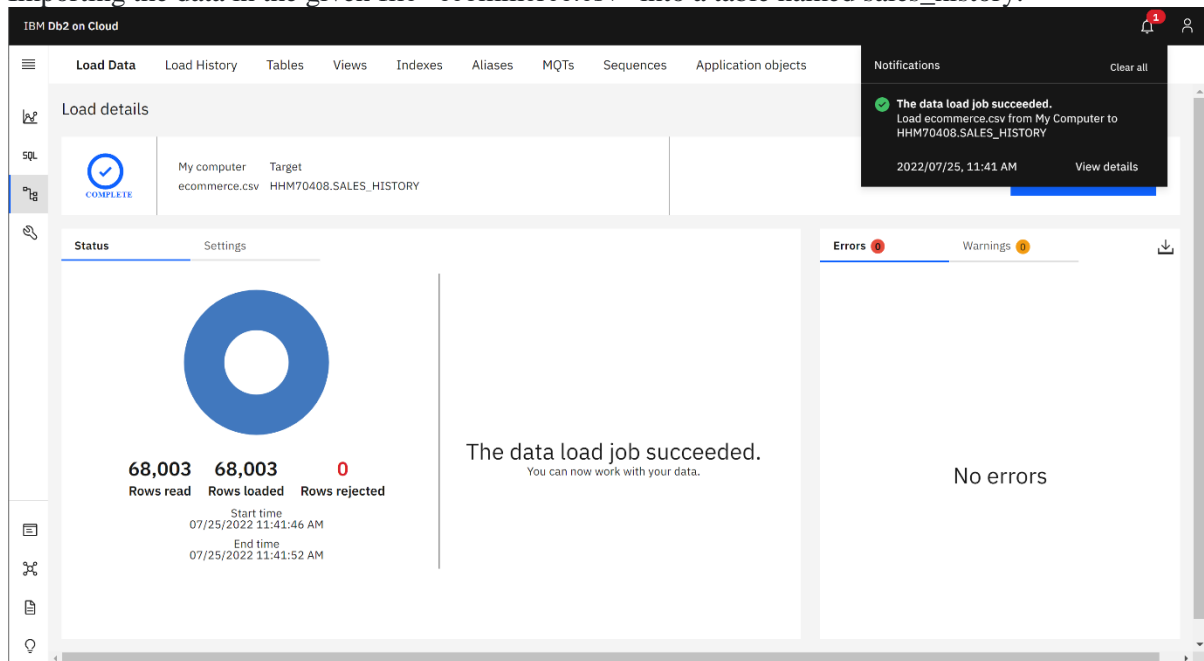
Module 4:

Scenario

You are a data engineer at an e-commerce company. Your company has finished setting up a data warehouse. Now you are assigned the responsibility to design a reporting dashboard that reflects the key metrics of the business.

Load data into the data warehouse (IBM DB2)

Importing the data in the given file “ecommerce.csv” into a table named sales_history.



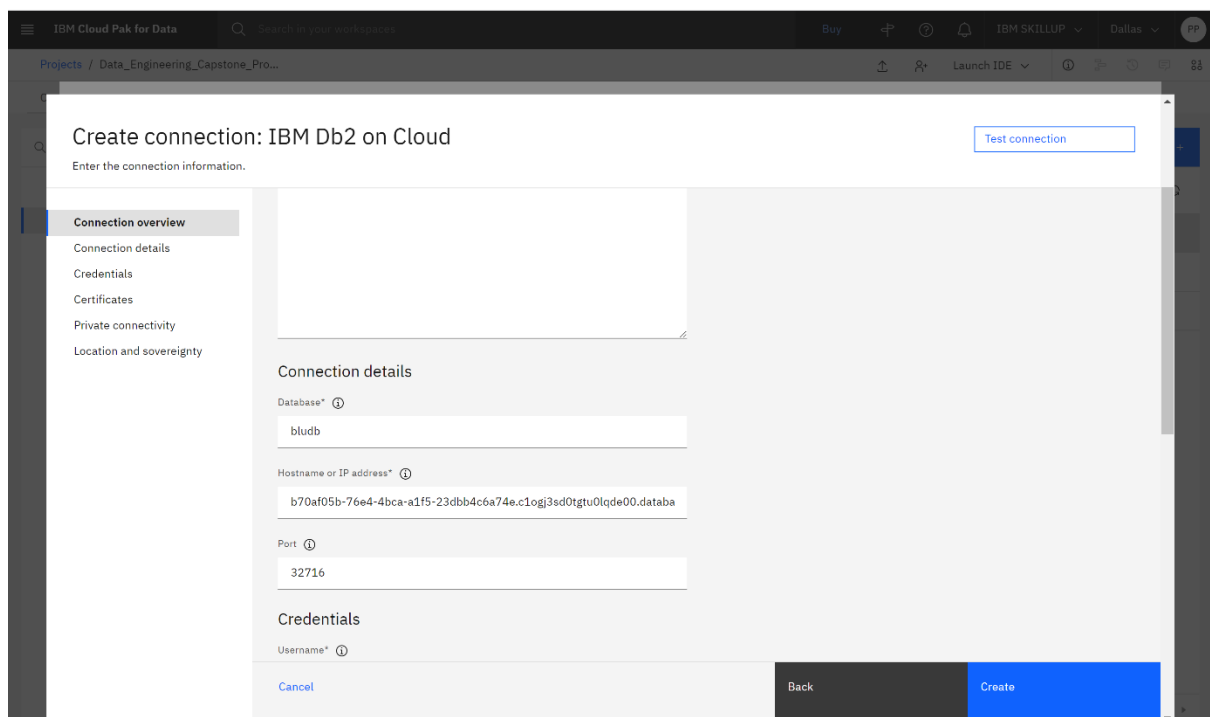
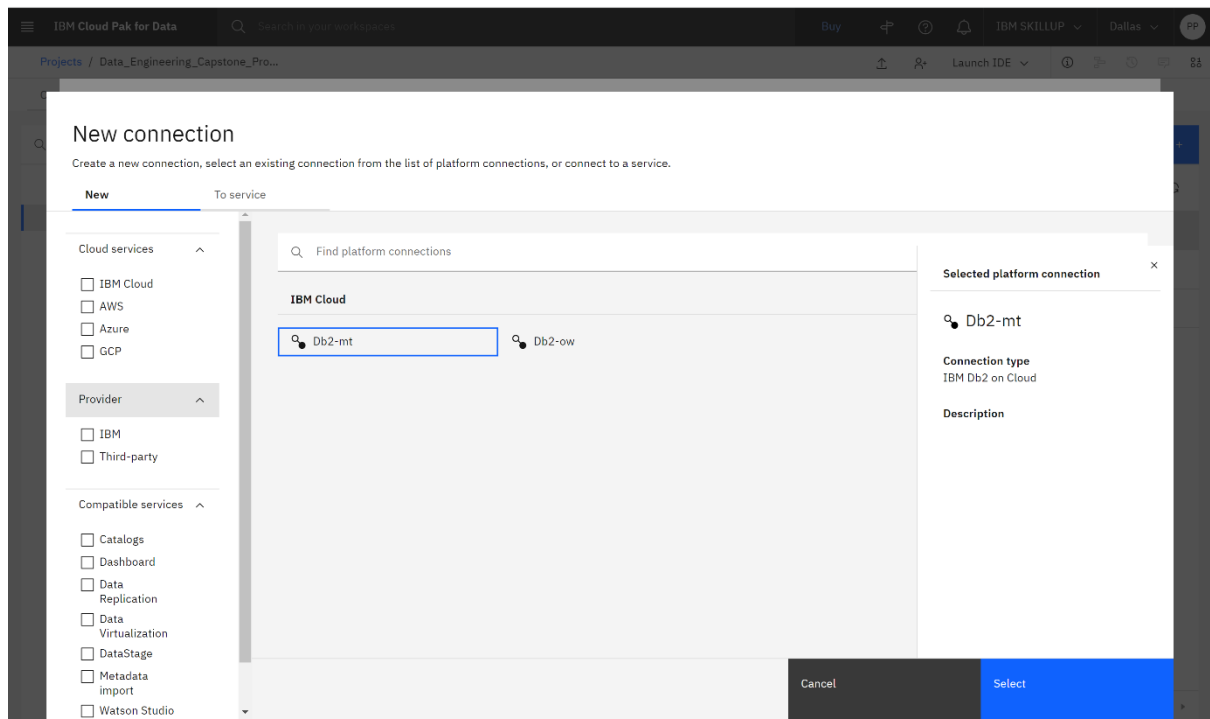
Listing the first 10 rows in the sales_history table to check the loaded data.

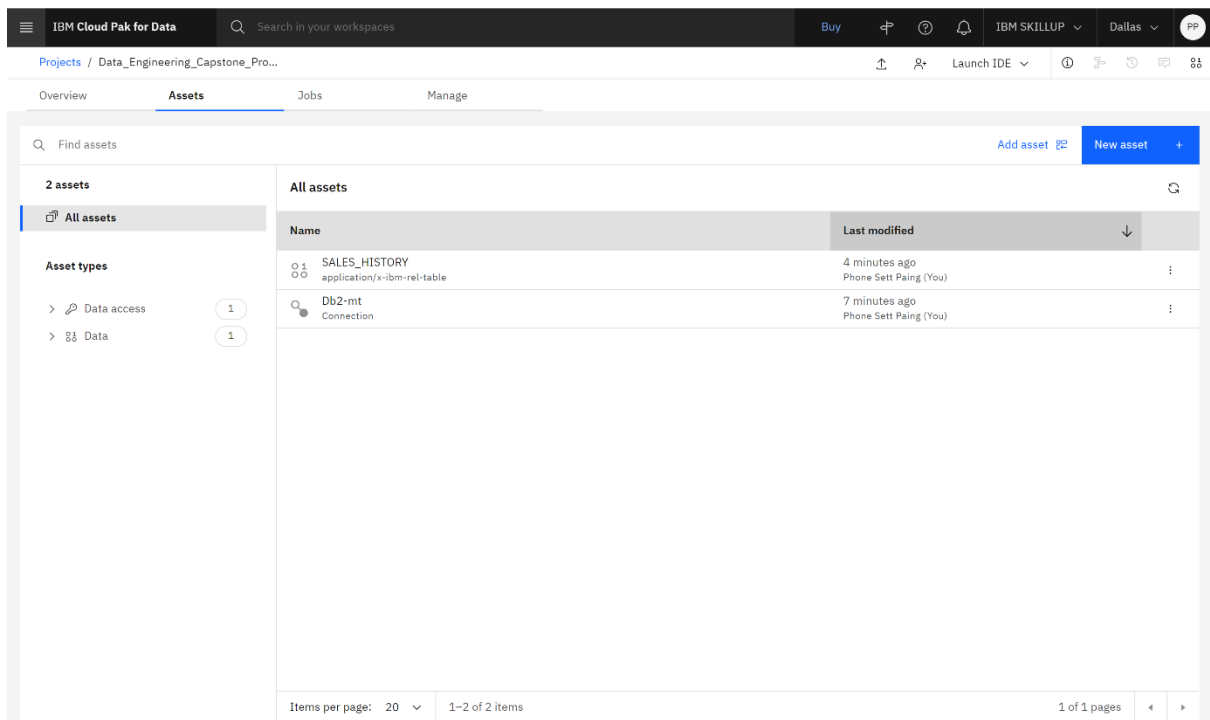
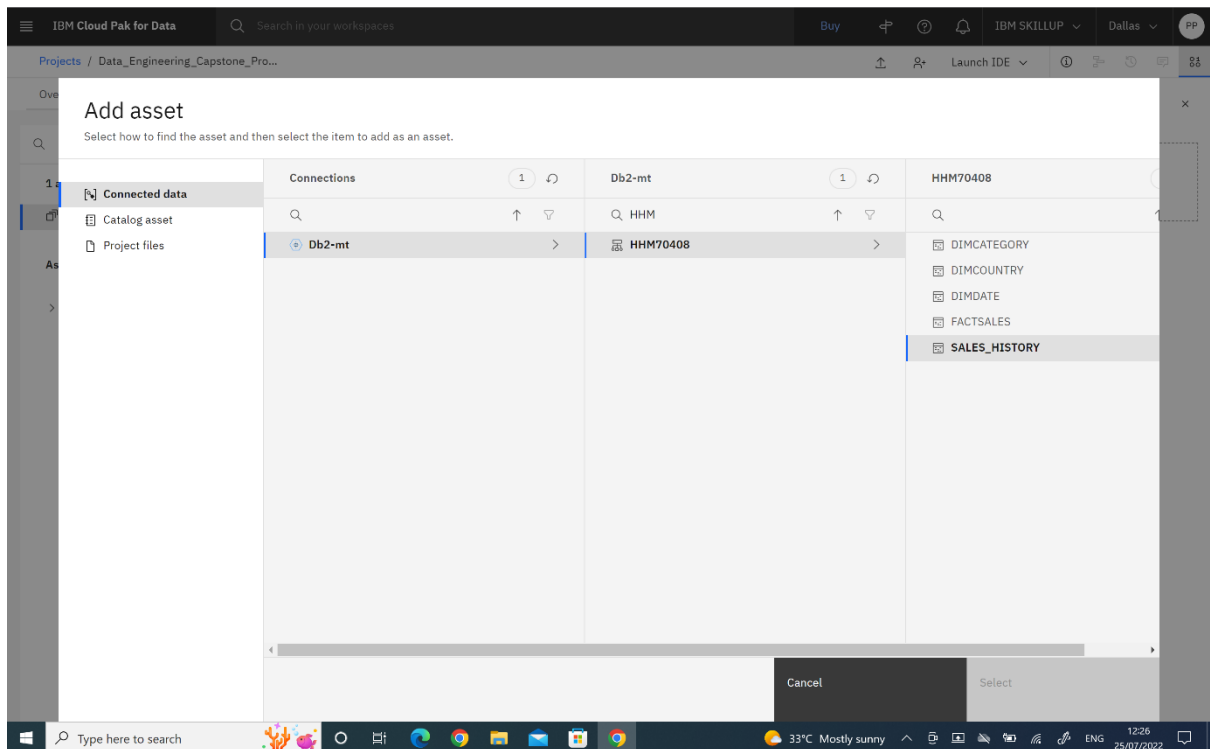
The screenshot shows the IBM Db2 on Cloud 'Run SQL' interface. The SQL query entered is 'SELECT * FROM sales_history LIMIT 10;'. The result set is displayed in a table with the following columns: DAY, MONTH, QUARTER, YEAR, CATEGORY, ITEM, and PRICE. The first 10 rows of data are shown.

DAY	MONTH	QUARTER	YEAR	CATEGORY	ITEM	PRICE
1	1	1	2020	Electronics	Tablet	1131
1	1	1	2020	Ebooks	Fiction	15
1	1	1	2020	Electronics	Laptop	866
1	1	1	2020	Electronics	Mobile	903
1	1	1	2020	Ebooks	Fiction	8
1	1	1	2020	Ebooks	Self Help	7
1	1	1	2020	Music	Pop	11
1	1	1	2020	Ebooks	Fiction	9
1	1	1	2020	Music	Pop	13
1	1	1	2020	Music	Pop	14

Create data source in IBM Cognos


Creating connection between sales_history table in IBM DB2 and Watson Studio with added Cognos Dashboard Embedded (CDE) service to use it as a data source.






Create dashboard using IBM Cognos

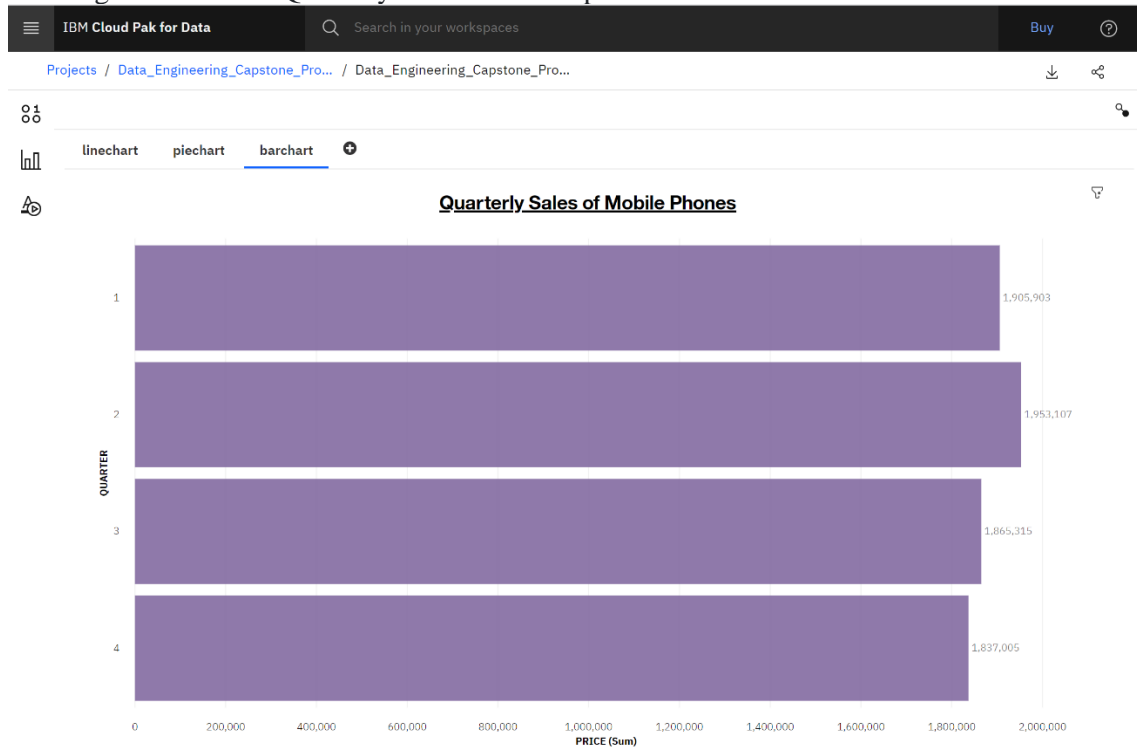
Creating a line chart of month wise total sales for the year 2020.



Creating a pie chart of category wise total sales.



Creating a bar chart of Quarterly sales of mobile phones.



Module 5.1:

Scenario

You are a data engineer at an e-commerce company. You need to keep data synchronized between different databases/data warehouses as a part of your daily routine. One task that is routinely performed is the sync up of staging data warehouse and production data warehouse. Automating this sync up will save you a lot of time and standardize your process. You will be given a set of python scripts to start with. You will use/modify them to perform the incremental data load from MySQL server which acts as a staging warehouse to the IBM DB2 which is a production data warehouse. This script will be scheduled by the data engineers to sync up the data between the staging and production data warehouse.

Automate loading of incremental data into the data warehouse

Creating the function `get_last_rowid()` that connect to the DB2 data warehouse and return the last rowid of the table `sales_data`.

```
automation.py > ...
36
37 # Find out the last rowid from DB2 data warehouse
38 # The function get_last_rowid must return the last rowid of the table sales_data of
39
40 def get_last_rowid():
41     SQL = "SELECT MAX(rowid) FROM SALES_DATA"
42     stmt = ibm_db.exec_immediate(conn, SQL)
43     tuple = ibm_db.fetch_tuple(stmt)
44     while tuple != False:
45         last_row = tuple[0]
46         return last_row
47
48 last_row_id = get_last_rowid()
49 print("\nLast row id on production datawarehouse = ", last_row_id)
50
```

Creating the function `get_latest_records()` that connect to MySQL database and return all records later than the given `last_rowid` from DB2 data warehouse.

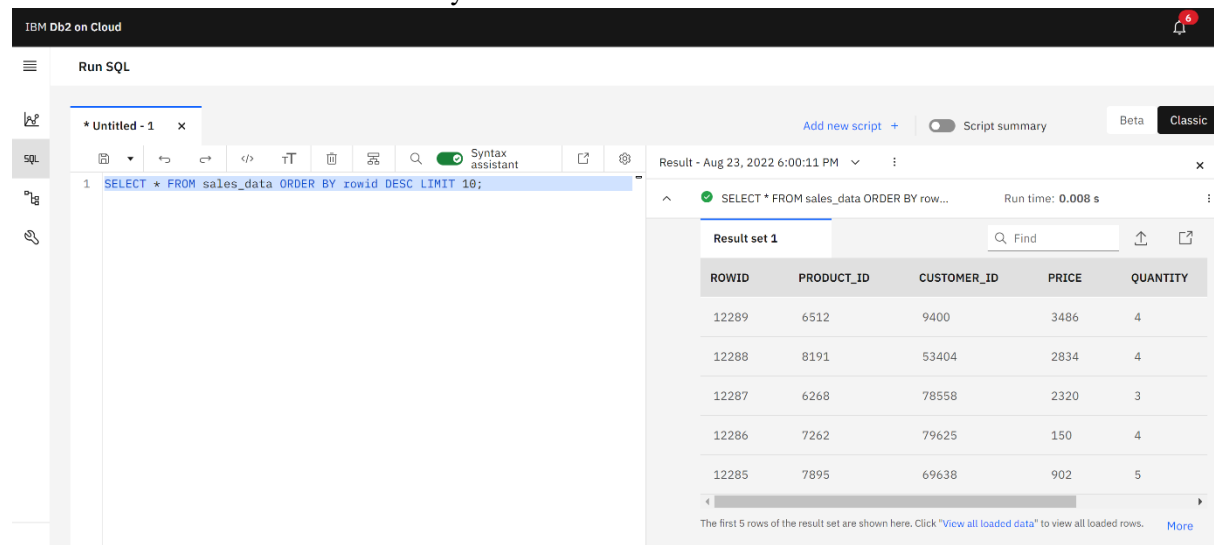
```
51 # List out all records in MySQL database with rowid greater than the one on the Data wa
52 # The function get_latest_records must return a list of all records that have a rowid g
53
54 def get_latest_records(rowid):
55     latest_row = []
56     SQL = f"SELECT * FROM sales_data WHERE rowid > {rowid}"
57     cursor.execute(SQL)
58     for row in cursor.fetchall():
59         latest_row.append(row)
60     return latest_row
61
62 new_records = get_latest_records(last_row_id)
63
64 print("\nNew rows on staging datawarehouse = ", len(new_records))
65
```

Creating the function `insert_records()` that connect to the DB2 data warehouse and insert all the given records (latest records from MySQL database).

```
65
66 # Insert the additional records from MySQL into DB2 data warehouse.
67 # The function insert_records must insert all the records passed to it into the sales_data table in IBM
68
69 def insert_records(records):
70     for record in records:
71         SQL = f"INSERT INTO SALES_DATA(rowid,product_id,customer_id,quantity) VALUES {record}"
72         ibm_db.exec_immediate(conn,SQL)
73
74 insert_records(new_records)
75 print("\nNew rows inserted into production datawarehouse = ", len(new_records))
76
77 # disconnect from mysql warehouse
78 connection.close()
79 # disconnect from DB2 data warehouse
80 ibm_db.close(conn)
81 # End of program
82
```

Testing the data synchronization

Data in DB2 data warehouse before synchronization

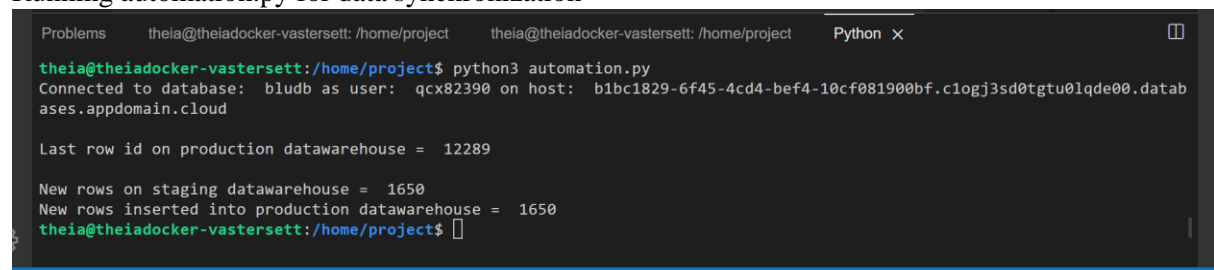


The screenshot shows the IBM Db2 on Cloud interface. The SQL editor contains the query: `SELECT * FROM sales_data ORDER BY rowid DESC LIMIT 10;`. The results pane displays the following data:

ROWID	PRODUCT_ID	CUSTOMER_ID	PRICE	QUANTITY
12289	6512	9400	3486	4
12288	8191	53404	2834	4
12287	6268	78558	2320	3
12286	7262	79625	150	4
12285	7895	69638	902	5

The interface also shows a 'Script summary' toggle, a 'Beta' label, and a 'Classic' button. The results pane includes a search bar and a 'View all loaded data' link.

Running automation.py for data synchronization

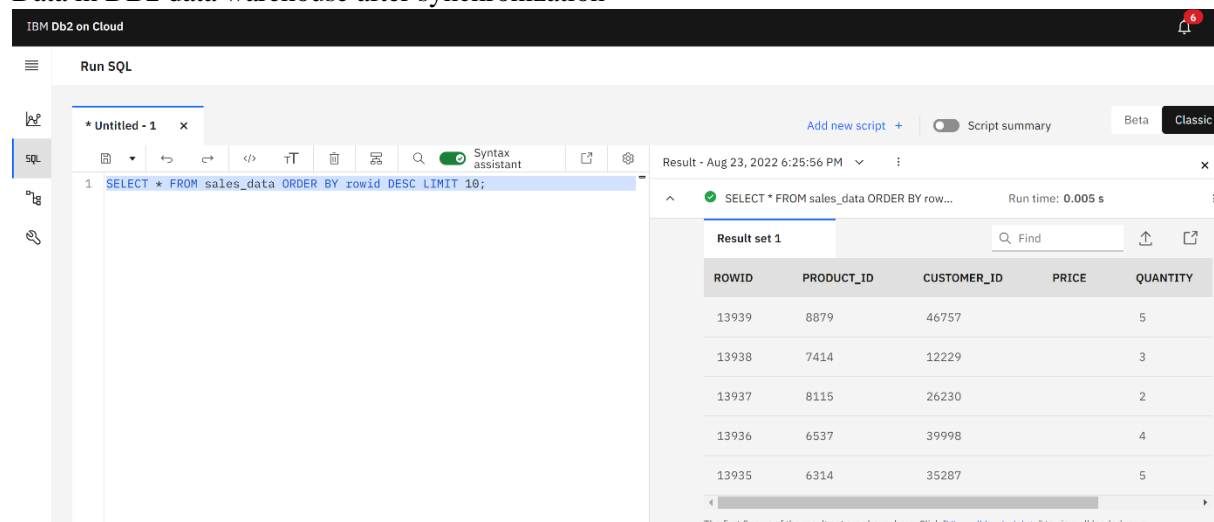


```
theia@theiadocker-vastersett: /home/project$ python3 automation.py
Connected to database: bludb as user: qcx82390 on host: b1bc1829-6f45-4cd4-bef4-10cf081900bf.clogj3sd0tgtu0lqde00.databases.appdomain.cloud

Last row id on production datawarehouse = 12289

New rows on staging datawarehouse = 1650
New rows inserted into production datawarehouse = 1650
theia@theiadocker-vastersett: /home/project$
```

Data in DB2 data warehouse after synchronization



The screenshot shows the IBM Db2 on Cloud interface after synchronization. The SQL editor contains the query: `SELECT * FROM sales_data ORDER BY rowid DESC LIMIT 10;`. The results pane displays the following data:

ROWID	PRODUCT_ID	CUSTOMER_ID	PRICE	QUANTITY
13939	8879	46757		5
13938	7414	12229		3
13937	8115	26230		2
13936	6537	39998		4
13935	6314	35287		5

The interface also shows a 'Script summary' toggle, a 'Beta' label, and a 'Classic' button. The results pane includes a search bar and a 'View all loaded data' link.

Data have been successfully synchronized.

Module 5.2:

Scenario

Write a pipeline that analyses the web server log file, extracts the required fields, transforms and load.

Creating Data Pipelines using Apache Airflow

Defining the DAG arguments.

```
process_web_log.py ×
airflow > dags > capstone > process_web_log.py > ...
1  # import the libraries
2  from datetime import timedelta
3  # The DAG object; we'll need this to instantiate a DAG
4  from airflow import DAG
5  # Operators; we need this to write tasks!
6  from airflow.operators.bash_operator import BashOperator
7  # This makes scheduling easy
8  from airflow.utils.dates import days_ago
9
10 #defining DAG arguments
11 # You can override them on a per-task basis during operator initialization
12 default_args = {
13     'owner': 'Phone Sett Paing',
14     'start_date': days_ago(0),
15     'email': ['vastersett@gmail.com'],
16     'email_on_failure': True,
17     'email_on_retry': True,
18     'retries': 1,
19     'retry_delay': timedelta(minutes=5),
20 }
21
22 # defining the DAG
```

Defining the DAG.

```
process_web_log.py ×
airflow > dags > capstone > process_web_log.py > ...
20 }
21
22 # defining the DAG
23 # define the DAG
24 dag = DAG(
25     dag_id='process_web_log',
26     default_args=default_args,
27     description='Analyzes the web server log file',
28     schedule_interval=timedelta(days=1),
29 )
30 |
31 # define the tasks
```

Creating a task to extract the ip address field from the web server log file into extracted_data.txt file.

```
process_web_log.py ●
airflow > dags > capstone > process_web_log.py > ...
29 )
30
31 # define the tasks
32 # define the first task (extract)
33 extract_data = BashOperator(
34     task_id='extract_data',
35     bash_command='cut -d" " -f1 /home/project/airflow/dags/capstone/accesslog.txt > extracted_data.txt',
36     dag=dag,
37 )
38
39 # define the second task (transform)
```

Creating a task to filter out all the occurrences of ip address “198.46.149.143” from extracted_data.txt file into transformed_data.txt file.

```
process_web_log.py ●
airflow > dags > capstone > process_web_log.py > ...
37 )
38
39 # define the second task (transform)
40 transform_data = BashOperator(
41     task_id='transform_data',
42     bash_command='grep -o "198.46.149.143" /home/project/airflow/dags/capstone/extracted_data.txt > transformed_data.txt',
43     dag=dag,
44 )
45
46 # define the third task (load)
```

Creating a task to archive the file transformed_data.txt file into weblog.tar file.

```
process_web_log.py ●
airflow > dags > capstone > process_web_log.py > ...
44 )
45
46 # define the third task (load)
47 load_data = BashOperator(
48     task_id='load_data',
49     bash_command='tar -cvf weblog.tar /home/project/airflow/dags/capstone/transformed_data.txt',
50     dag=dag,
51 )
52
53 # task pipeline
```

Defining the task pipeline.

```
process_web_log.py ●
airflow > dags > capstone > process_web_log.py > ...
51 )
52
53 # task pipeline
54 extract_data >> transform_data >> load_data
55 |
```

Getting the DAG operational

Submitting the DAG


```
Problems theia@theiadocker-vastersett: /home/project/airflow/dags/capstone X
theia@theiadocker-vastersett:/home/project/airflow/dags/capstone$ cp process_web_log.py $AIRFLOW_HOME/dags
theia@theiadocker-vastersett:/home/project/airflow/dags/capstone$
```

Unpausing the DAG

```
theia@theiadocker-vastersett:/home/project/airflow/dags/capstone$ airflow dags unpause process_web_log
Dag: process_web_log, paused: False
theia@theiadocker-vastersett:/home/project/airflow/dags/capstone$
```

Monitoring the DAG

← → ↻ 🏠 vastersett-8080.theiadocker-0-labs-prod-theiak8s-4-tor01.proxy.cognitiveclass.ai/home

 Airflow Security Browse Admin Docs

<input type="checkbox"/>	example_xcom_args example	airflow	<div><div></div><div></div><div></div><div></div></div>	None	
<input type="checkbox"/>	example_xcom_args_with_operators example	airflow	<div><div></div><div></div><div></div><div></div></div>	None	
<input type="checkbox"/>	latest_only example2 example3	airflow	<div><div></div><div></div><div></div><div></div></div>	4:00:00	2022-08-07, 10:01:21 ⓘ
<input type="checkbox"/>	latest_only_with_trigger example3	airflow	<div><div></div><div></div><div></div><div></div></div>	4:00:00	2022-08-07, 10:01:20 ⓘ
<input checked="" type="checkbox"/>	process_web_log	Phone Sett Paing	<div><div></div><div></div><div></div><div></div></div>	1 day, 0:00:00	2022-08-07, 00:00:00 ⓘ
<input type="checkbox"/>	test_utils example	airflow	<div><div></div><div></div><div></div><div></div></div>	None	
<input type="checkbox"/>	tutorial example	airflow	<div><div></div><div></div><div></div><div></div></div>	1 day, 0:00:00	2022-08-06, 14:01:24 ⓘ
<input type="checkbox"/>	tutorial_etl_dag example	airflow	<div><div></div><div></div><div></div><div></div></div>	None	
<input type="checkbox"/>	tutorial_taskflow_api_etl example	airflow	<div><div></div><div></div><div></div><div></div></div>	None	
<input type="checkbox"/>	tutorial_taskflow_api_etl_virtualenv example	airflow	<div><div></div><div></div><div></div><div></div></div>	None	

Module 6:

Analyse search terms on the e-commerce web server

Download the search term data set for the e-commerce web server and analyse the data set, load the sales forecast model and predict the sales for the year 2023.

Printing the number of rows and columns in the data set.

```
[7]: # Load the csv into a spark dataframe

[8]: df = spark.read.csv('searchterms.csv',header=True).cache()

[9]: # Print the number of rows and columns
     # Take a screenshot of the code and name it as shape.jpg)

[10]: print((df.count(),len(df.columns)))

[Stage 1:> (0 + 1) / 1]
(10000, 4)
```

Printing the top 5 rows in the data set.

```
[11]: # Print the top 5 rows
     # Take a screenshot of the code and name it as top5rows.jpg)

[12]: df.show(5)

+---+-----+---+-----+
|day|month|year|  searchterm|
+---+-----+---+-----+
| 12|   11|2021| mobile 6 inch|
| 12|   11|2021| mobile latest|
| 12|   11|2021|  tablet wifi|
| 12|   11|2021|laptop 14 inch|
| 12|   11|2021|   mobile 5g|
+---+-----+---+-----+
only showing top 5 rows
```

Finding the data type of the column “searchterm”.

```
[13]: # Find out the datatype of the column searchterm?
     # Take a screenshot of the code and name it as datatype.jpg)

[14]: df.printSchema()

root
 |-- day: string (nullable = true)
 |-- month: string (nullable = true)
 |-- year: string (nullable = true)
 |-- searchterm: string (nullable = true)
```

Finding the number of times, the term “gaming laptop” was searched.

```
[15]: # How many times was the term `gaming laptop` searched?
      # Take a screenshot of the code and name it as gaminglaptop.jpg

[16]: df.groupby("searchterm").count().filter("searchterm='gaming laptop'").show()

[Stage 13:=====> (58 + 8) / 75]
+-----+-----+
| searchterm|count|
+-----+-----+
| gaming laptop| 499|
+-----+-----+
```

Printing the top 5 frequently used search terms.

```
[17]: # Print the top 5 most frequently used search terms?
      # Take a screenshot of the code and name it as top5terms.jpg

[18]: df.groupby("searchterm").count().sort('count',ascending=False).show(5)

[Stage 15:=====> (176 + 11) / 200]
+-----+-----+
| searchterm|count|
+-----+-----+
| mobile 6 inch| 2312|
| mobile 5g| 2301|
| mobile latest| 1327|
| laptop| 935|
| tablet wifi| 896|
+-----+-----+
only showing top 5 rows
```

Loading the sales forecast model.

```
sales_prediction.model/data/.part-00000-1db9fe2f-4d93-4b1f-966b-3b09e72d664e-
sales_prediction.model/data/._SUCCESS.crc
```

```
[22]: # Load the sales forecast model.
      # Take a screenshot of the code and name it as Loadmodel.jpg
      from pyspark.ml.regression import LinearRegressionModel

      model=LinearRegressionModel.load('sales_prediction.model')

[23]: # Using the sales forecast model, predict the sales for the year of 2023.
      # Take a screenshot of the code and name it as forecast ing
```

Predicting the sales for the year 2023, using the sales forecast model.

```
[23]: # Using the sales forecast model, predict the sales for the year of 2023.  
# Take a screenshot of the code and name it as forecast.jpg
```

```
[24]: from pyspark.ml.feature import VectorAssembler  
  
def predict(year):  
    assembler = VectorAssembler(inputCols=["year"],outputCol="features")  
    data=[[year,0]]  
    columns=["year","sales"]  
    _ = spark.createDataFrame(data,columns)  
    __ = assembler.transform(_).select('features','sales')  
    predictions = model.transform(__)  
    predictions.select('prediction').show()  
  
predict(2023)
```

```
+-----+  
|      prediction|  
+-----+  
|175.16564294006457|  
+-----+
```